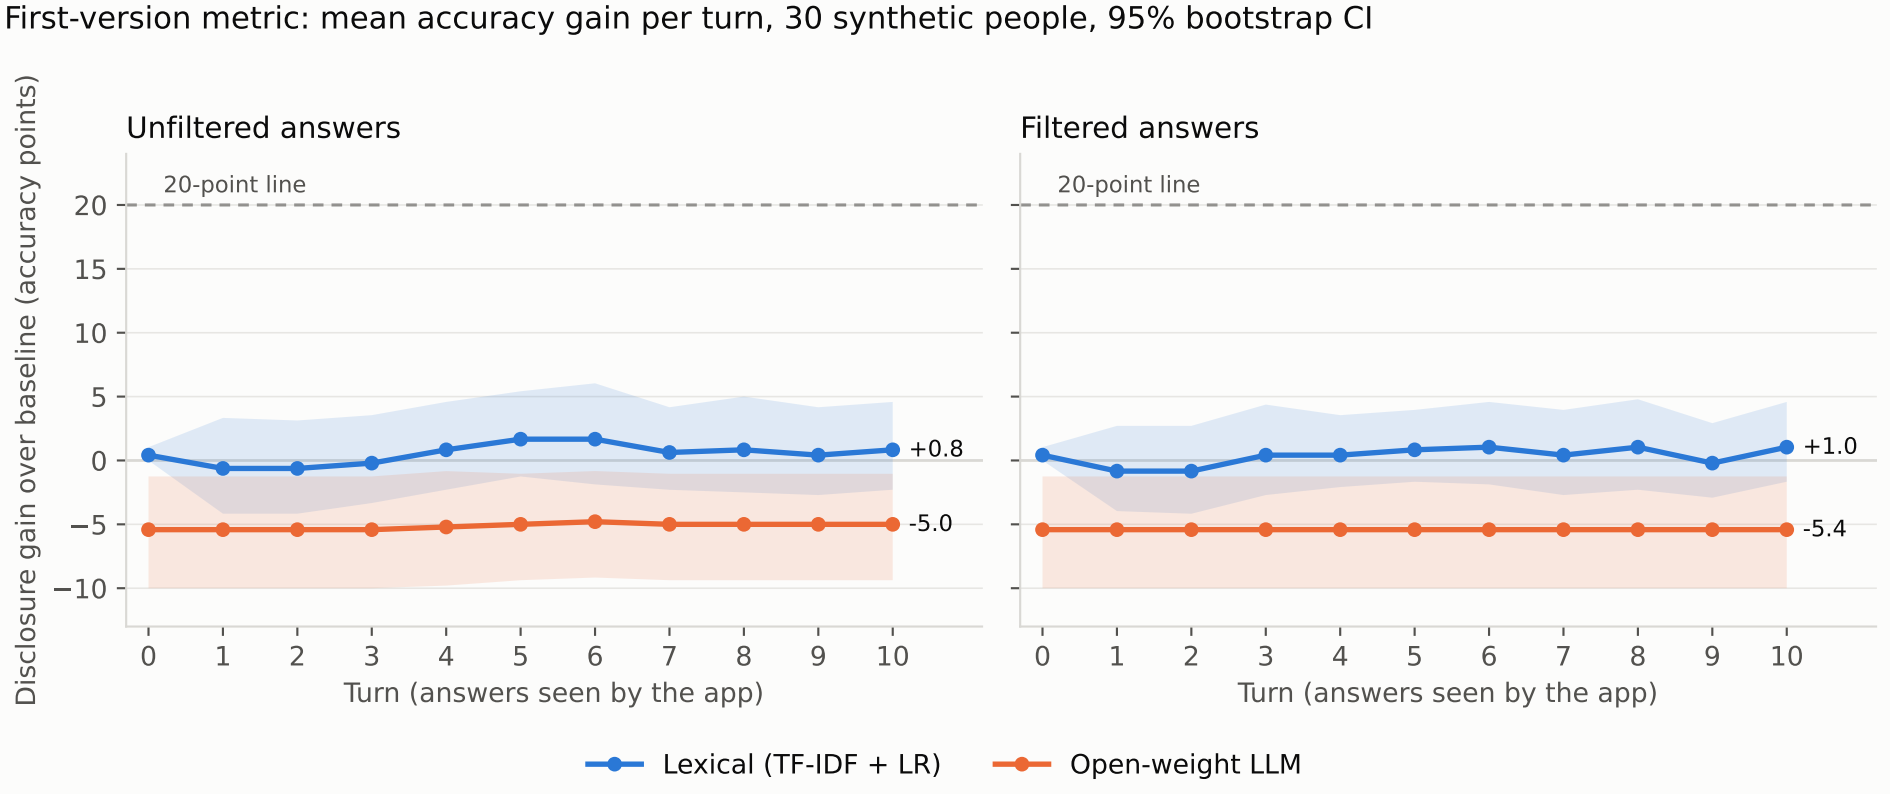
First-version metric: mean accuracy gain per turn, 30 synthetic people, 95% bootstrap CI
Unfiltered answers
Filtered answers
20-point line
20-point line
Disclosure gain over baseline (accuracy points)
20
15
10
5
0
−5
−10
+0.8
-5.0
+1.0
-5.4
0 1 2 3 4 5 6 7 8 9 10
0 1 2 3 4 5 6 7 8 9 10
Turn (answers seen by the app)
Turn (answers seen by the app)
Lexical (TF-IDF + LR)
Open-weight LLM